

# Planning, Learning and Decision Making:

## Homework 2. Markov decision problems

Mykhaylo Marfeychuk, Roy De Prins

16 October 2018

### 1 Exercise 1) a

State Space  $\mathcal{X}$  =

{TopLeft, TopRight, BottomLeft, BottomRight,  
Red, TopLeftRed, TopRightRed, BottomLeftRed, BottomRightRed,  
Blue, TopLeftBlue, TopRightBlue, BottomLeftBlue, BottomRightBlue,  
Goal }.

Action Space  $\mathcal{A}$  =

{Up, Down, Left, Right }.

### 2 Exercise 1) b

Transition Probability Matrix for action 'right':

|      | TL  | TR  | BL  | BR  | Red | TLR | TRR | BLR | BRR | Blue | TLB | TRB | BLB | BRB | Goal |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|------|-----|-----|-----|-----|------|
| TL   | 0.2 | 0.8 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| TR   | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| BL   | 0.0 | 0.0 | 0.2 | 0.8 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| BR   | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| Red  | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| TLR  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.2 | 0.8 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| TRR  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| BLR  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.2 | 0.8 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| BRR  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  |
| Blue | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.2  | 0.0 | 0.8 | 0.0 | 0.0 | 0.0  |
| TLB  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.2 | 0.8 | 0.0 | 0.0 | 0.0  |
| TRB  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 1.0 | 0.0 | 0.0 | 0.0  |
| BLB  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.2 | 0.8 | 0.0  |
| BRB  | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.2 | 0.8  |
| Goal | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 | 0.0 | 1.0  |

Cost Function for action 'right':

|      | cost_right |
|------|------------|
| TL   | 0.5        |
| TR   | 1.0        |
| BL   | 0.5        |
| BR   | 1.0        |
| Red  | 1.0        |
| TLR  | 0.5        |
| TRR  | 1.0        |
| BLR  | 0.5        |
| BRR  | 1.0        |
| Blue | 0.0        |
| TLB  | 0.0        |
| TRB  | 1.0        |
| BLB  | 0.0        |
| BRB  | 0.0        |
| Goal | 0.0        |

### 3 Excercise 1) c

We compute the cost-to-go function with the 'Value Iteration, 2.0' algorithm. So by computing formula 1:

$$\mathbf{J}^t = \mathbf{c}_{right} + \gamma \mathbf{P}_{right} \mathbf{J}^{t-1} \quad (1)$$

The cost-to-go function associated with the policy 'right' is:

| cost-to-go |
|------------|
| 9.39       |
| 10.0       |
| 9.39       |
| 10.0       |
| 10.0       |
| 9.39       |
| 10.0       |
| 9.39       |
| 10.0       |
| 8.78       |
| 8.78       |
| 10.0       |
| 0.0        |
| 0.0        |
| 0.0        |

This calculation is done by the program (in Python) in Figure 1.

```

1  import numpy as np
2
3  cost_right = np.array([[0.5], [1], [0.5], [1], [1], [0.5], [1], [0.5], [1], [0], [0], [1], [0], [0], [0]])
4  prob_right = np.array([[0.2, 0.8, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
5                          [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
6                          [0.0, 0.0, 0.2, 0.8, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
7                          [0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
8                          [0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
9                          [0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
10                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
11                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
12                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0],
13                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8, 0.0, 0.0, 0.0, 0.0],
14                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8, 0.0, 0.0, 0.0],
15                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0],
16                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8, 0.0],
17                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.2, 0.8],
18                         [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 1.0]])
19
20  gamma = 0.9
21
22  J = np.zeros((15,1))
23  err = 1
24  i = 0
25
26  while err > 1e-15:
27      Jnew = cost_right + gamma * prob_right.dot(J)
28      err = np.linalg.norm(Jnew - J)
29      i += 1
30      J = Jnew
31
32  print("The cost-to-go function is \n", J)

```

Run: HW2 x

```

The cost-to-go function is
[[ 9.3902439]
 [10.      ]
 [ 9.3902439]
 [10.      ]
 [10.      ]
 [ 9.3902439]
 [10.      ]
 [ 9.3902439]
 [10.      ]
 [ 8.7804878]
 [ 8.7804878]
 [10.      ]
 [ 0.      ]
 [ 0.      ]
 [ 0.      ]]

```

Figure 1: The Python program for calculating to cost-to-go function.